

The AI Governance Checklist for C-Suite

12 Questions Every Board Must Answer in 2026

Published by AI Times — Editorial signals for C-suite on AI, governance, and infrastructure.

Why this checklist

In our editorial view at AI Times: **95% of enterprise AI projects deliver zero measurable ROI** (Hastings Journal, MIT research, December 2025). The reason isn't the technology — it's the absence of governance discipline at board level. Most C-suites are approving AI initiatives without the same rigor they apply to financial controls, M&A, or cybersecurity.

This checklist is a 12-question instrument. Run it through your AI initiative — internal copilots, customer-facing agents, third-party model integrations, board-level AI advisory. If you can't answer at least 9 of 12 with concrete documentation, your governance is theatre.

Section A — Strategic Mandate (3 questions)

1. Who is accountable on the board for AI strategy outcomes?

- Not "the CTO." A named director with explicit P&L accountability and quarterly review obligation.
- Document the delegation chain: board → committee → exec → working group → engineering.
- Reference: McKinsey 2025 "State of AI" — only 28% of enterprises have explicit board accountability.

2. What is the failure mode tolerance — by use case, not aggregated?

- Customer-facing AI: hallucination tolerance, regulatory exposure, brand-damage downside.
- Internal tooling: productivity-loss downside, employee-trust erosion, shadow-IT proliferation.
- Trading / quant AI: tail-risk magnitude under model drift.
- Different use cases require different governance — one policy fits none.

3. What is the explicit kill-switch protocol?

- Who can pause an AI system, by what process, in what timeframe (real-time / 1 hour / 24 hours)?
- What contractual obligations cascade if a third-party AI provider experiences an incident?
- Document the decision tree before you need it.

Section B — Data Foundations (3 questions)

4. Data lineage — can you audit every training input to a model that influences a customer decision?

- "We use OpenAI / Anthropic / etc." is not an answer — that pushes the question to vendor agreements.
- Map the data flows: customer data → prompts → models → outputs → downstream decisions.
- Required for EU AI Act compliance, increasingly for state-level US regulation.

5. PII / sensitive-data handling in prompts — who reviewed your prompt corpus?

- Most enterprise prompt injection incidents in 2025 originated from prompts containing PII / financial data / customer secrets being logged by AI providers.
- Audit checklist: prompt logs, vendor data-residency claims, third-party access (sub-processors).

6. Model selection — is there a documented decision rationale per AI use case?

- Why GPT-5 vs Claude vs Gemini vs Llama 3 self-hosted?
- Cost, latency, capability, data-residency, brand-alignment, lock-in risk.
- "Our team likes Anthropic" is not a board-defensible answer.

Section C — Operational Controls (3 questions)

7. Hallucination detection — what is your false-positive rate target and current measurement?

- Define use-case specific tolerances (legal advice 0% / marketing copy 5% / internal Q&A 2%).
- Establish a sample-and-grade pipeline with human review of model outputs.
- Track drift over time.

8. Prompt injection / jailbreak resistance — when was your last red-team test?

- Industry baseline: quarterly red-team for customer-facing AI, semi-annual for internal.
- Vendor claims of "guardrails" are insufficient — assume adversarial users.

9. AI vendor concentration — what is your operational dependency on a single LLM provider?

- If OpenAI / Anthropic / Google has a 4-hour outage, what fails in your stack?
- Multi-provider fallback architecture is the new operational resilience standard.
- Reference: AWS Bedrock / Azure Foundry multi-model strategies.

Section D — Compliance & Regulatory (3 questions)

10. EU AI Act classification — have you mapped your AI systems to risk tiers?

- Unacceptable / High-risk / Limited / Minimal — each tier carries distinct documentation obligations.
- Deadline: full enforcement phases through 2026-2027.
- Reference: EU Regulation 2024/1689, Articles 6, 9, 13, 14.

11. US state-level AI regulation — Colorado SB 205, California AB 2930, NY State CIDC Act?

- Patchwork of state-level rules with conflicting requirements (transparency / audit / employer obligations).
- Document jurisdictional exposure: where customers / employees / data are located.

12. Disclosure to customers — when do you tell users they're interacting with AI?

- EU AI Act Article 50 requires it for chatbots and emotion-recognition systems.
- Several US states + FTC guidance moving the same direction.
- Document your disclosure standard, audit recent customer interactions for compliance.

Scoring

- **9-12 ✓ answered with documentation** — governance discipline at board-grade. Continue quarterly review.
- **5-8 ✓** — operational gaps. Schedule 90-day remediation with named owners.
- **<5 ✓** — governance theatre. Halt new AI deployments until baseline is in place.

What to do next

1. **Audit your AI inventory** — list every AI system currently in production or piloting.
2. **Apply this checklist** — score each system, document gaps.
3. **Brief the board** — quarterly AI governance review on the agenda.
4. **Establish kill-switch protocols** — before you need them.

References

- EU Regulation 2024/1689 (EU AI Act) — eur-lex.europa.eu
- Hastings Journal, "Why 95% of enterprise AI projects fail to deliver ROI," December 2025
- MIT Sloan, "AI Governance for the Board," October 2025
- McKinsey, "The state of AI in 2025"
- Boston Consulting Group, "AI Adoption 2024," October 2024
- Colorado SB 205 (2024), California AB 2930 (2024), NY State CIDC Act

About AI Times

We publish editorial signals on AI policy, governance, and infrastructure for C-suite executives. Daily reels on Instagram. Weekly deep-dives on Beehiiv.

Follow @aittimes on Instagram for daily AI & Crypto signals.

Subscribe to the AI Times Brief — weekly editorial newsletter for C-suite, no fluff: [Beehiiv signup — pending]

Editorial inquiries: ttimesai@gmail.com

© 2026 AI Times. Published 13 May 2026. Lead magnet — share freely with attribution.